

## CO527: Advanced Database Systems

### Lab 03 - Query Optimization

Ranage R.D.P.R.  
E/19/310

1. Use explain to analyze the outputs of following two simple queries which use only one table access.

I. **SELECT \* FROM departments WHERE deptname = 'Finance';**

```
MySQL localhost:3306 ssl company SQL > EXPLAIN SELECT * FROM department WHERE dept_name = 'Finance';
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table      | partitions | type | possible_keys | key  | key_len | ref | rows | filtered | Extra           |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1  | SIMPLE      | department | NULL       | ALL  | NULL          | NULL | NULL    | NULL | 9    | 11.11    | Using where     |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.0013 sec)
```

II. **SELECT \* FROM departments WHERE deptno = 'd002';**

```
MySQL localhost:3306 ssl company SQL > EXPLAIN SELECT * FROM department WHERE dept_no = 'd002';
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| id | select_type | table      | partitions | type | possible_keys | key  | key_len | ref | rows | filtered | Extra           |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
| 1  | SIMPLE      | department | NULL       | const | PRIMARY       | PRIMARY | 16      | const | 1    | 100      | NULL            |
+----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
1 row in set, 1 warning (0.0008 sec)
```

What conclusions you can draw from the results?

1. Query 1 (dept\_name = 'Finance'):

The query uses a simple table scan (type: ALL), meaning it examines all rows in the departments table.

There are no possible keys or indexes that can be utilized for this query (possible\_keys: NULL, key: NULL).

The WHERE clause is applied directly to the table (Using where).

The number of rows examined is 9.

Conclusion: This query is not utilizing any indexes, resulting in a full table scan, and could potentially benefit from an index on the dept\_name column to improve performance.

2. Query 2 (dept\_no = 'd002'):

The query uses a const access method, which indicates a single-row lookup based on a constant value.

The PRIMARY key is used for this query (possible\_keys: PRIMARY, key: PRIMARY), and the lookup is based on the primary key index.

The key length is 16 bytes, matching the length of the primary key.

Only one row is examined.

Conclusion: This query benefits from the presence of the primary key index on the dept\_no column, resulting in efficient single-row lookup.

2. Start by creating the initial tables emplist and titleperiod as follows. These derived tables need to contain only the columns involved in the query.

I. **create table emplist select emp\_no, first\_name from employees;**

II. **create table titleperiod select emp\_no, title, datediff(to\_date, from\_date) as period FROM titles;**

```
MySQL localhost:3306 ssl company SQL > create table emplist select emp_no, first_name from employee;
Query OK, 300024 rows affected (4.0564 sec)

Records: 300024 Duplicates: 0 Warnings: 0
MySQL localhost:3306 ssl company SQL > create table titleperiod select emp_no, title, datediff(to_date, from_date) as period FROM title;
Query OK, 443306 rows affected (6.1473 sec)

Records: 443306 Duplicates: 0 Warnings: 0
```

Now write the query that gives the desired information in the required format.

```
Records: 443306 Duplicates: 0 Warnings: 0
MySQL localhost:3306 ssl company SQL> select e.first_name, t.period
-> from emplist e JOIN titleperiod t ON e.emp_no = t.emp_no
-> WHERE t.period > 4000
-> LIMIT 20;
```

first_name	period
Georgi	2926512
Bezalel	2922821
Parto	2923065
Chirstian	2923067
Kyoichi	2922781
Anneke	2925011

Analyze the output of applying EXPLAIN to the above query explaining each value. Note that the tables are in their initial unindexed state.

```
MySQL localhost:3306 ssl company SQL> EXPLAIN select e.first_name, t.period from emplist e JOIN titleperiod t ON e.emp_no = t.emp_no WHERE t.period > 4000 LIMIT 20;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	t	NULL	ALL	NULL	NULL	NULL	NULL	442668	33.33	Using where
1	SIMPLE	e	NULL	ALL	NULL	NULL	NULL	NULL	299751	10	Using where; Using join buffer (hash join)

2 rows in set, 1 warning (0.0008 sec)

The query that involves two tables, t and e, and its execution process is explained in the EXPLAIN result. A full table scan of both tables is performed; this is indicated by the type: ALL, indicating that each table's rows are looked over. Moreover, the tables are in their original unindexed state because there are no potential keys and no indexes are used. As a result, the query only uses sequential scans of all the tables, which can result in less than ideal speed, especially when dealing with big datasets. Furthermore, the WHERE clause is used in both tables, suggesting that filtering is done while the query is being executed. The fact that Table e uses a join buffer for a hash join operation is noteworthy since it implies that the join between the two tables is processed by a hash join algorithm. Nevertheless, since the optimizer expects a large number of rows to be analyzed, the absence of indexes and the reliance on complete table scans suggest possible performance bottlenecks. Thus, it's best to construct the right indexes on the columns that are used in join conditions and filtering predicates in order to improve query performance. Queries can be executed more quickly and with greater efficiency when pertinent columns are indexed. This allows the query optimizer to make use of more effective access techniques, including index scans.

**What could be the number of row combinations that MySQL would need to check?**

No. of row combinations = titleperiod table rows x emplist table rows  
= 442929 x 299715  
= 132,752,465,235

3.

I. Create indexes on the columns used to join the tables. In the emplist table, emp\_no can be used as a primary key because it uniquely identifies each row.

```
MySQL localhost:3306 ssl company SQL> alter table emplist
-> Add primary key(emp_no);
Query OK, 0 rows affected (1.5009 sec)
```

```
MySQL localhost:3306 ssl company SQL> alter table titleperiod
-> add index(emp_no);
Query OK, 0 rows affected (1.3945 sec)
Records: 0 Duplicates: 0 Warnings: 0
```

```
MySQL localhost:3306 ssl company SQL> EXPLAIN select e.first_name, t.period from emplist e JOIN titleperiod t ON e.emp_no = t.emp_no WHERE t.period > 4
000 LIMIT 20;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	e	NULL	ALL	PRIMARY	NULL	NULL	NULL	2995045	100	NULL
1	SIMPLE	t	NULL	ref	emp_no	emp_no	5	company.e.emp_no	1	33.33	Using where

2 rows in set, 1 warning (0.0034 sec)

**Optimizing Table 'e' Access:** Ensure that the query filters rows based on the primary key of table 'e' to utilize the primary key index efficiently. If the primary key is not used for filtering, consider whether it is necessary for the query. If not, the primary key index can potentially be dropped to save space and improve write performance.

## Query Rewriting Techniques

1. **USE/FORCE INDEX** - To force MySQL to use an index.
2. **IGNORE INDEX** - To tell MySQL to ignore an index.

```
MySQL localhost:3306 ssl company SQL > EXPLAIN select first_name
-> from emplist
-> FORCE INDEX(PRIMARY)
-> where emp_no > 1000;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	emplist	NULL	range	PRIMARY	PRIMARY	4	NULL	1	100	Using where

```
1 row in set, 1 warning (0.0085 sec)
```

```
MySQL localhost:3306 ssl company SQL> EXPLAIN SELECT STRAIGHT_JOIN first_name from emplist
-> FORCE INDEX(PRIMARY)
-> WHERE emp_no>1000;
```

id	select_type	table	partitions	type	possible_keys	key	key_len	ref	rows	filtered	Extra
1	SIMPLE	emplist	NULL	range	PRIMARY	PRIMARY	4	NULL	1	100	Using where

```
1 row in set. 1 warning (0.0014 sec)
```

Therefore, based on the information provided, we can conclude that MySQL's EXPLAIN function is a strong tool for SQL query optimization. It offers details on how a specific query is carried out by MySQL.

- Understanding the execution of queries
- locating bottlenecks
- making use of indexes
- joining operations efficiently
- optimizing select queries.